



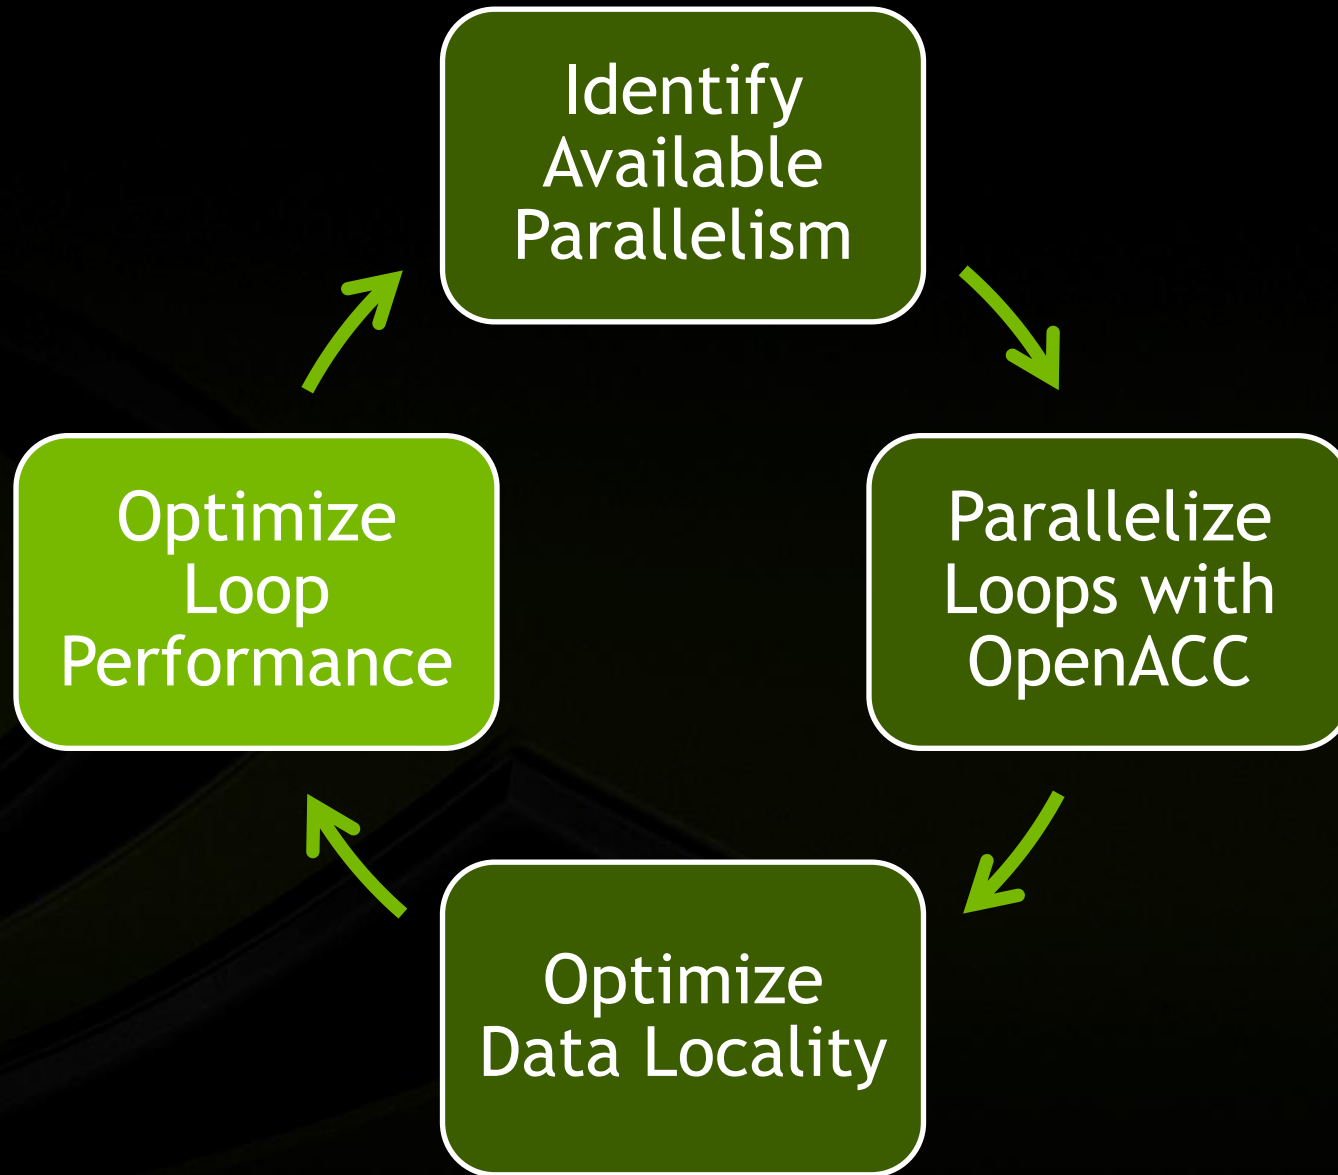
ADVANCED OPENACC PROGRAMMING

DR. CHRISTOPH ANGERER, NVIDIA

*) THANKS TO JEFF LARKIN, NVIDIA, FOR THE SLIDES

AGENDA

- ▶ Optimizing OpenACC Loops
- ▶ Routines
- ▶ Update Directive
- ▶ Asynchronous Programming
- ▶ Multi-GPU Programming
- ▶ OpenACC Interoperability
- ▶ Atomic Directive
- ▶ Misc. Advice & Techniques
- ▶ Q&A



SPARSE MATRIX/VECTOR PRODUCT

```
99      do i=1,a%num_rows
100          tmpsum = 0.0d0
101          row_start = arow_offsets(i)
102          row_end   = arow_offsets(i+1)-1
103          do j=row_start,row_end
104              acol = acols(j)
105              acoef = acoefs(j)
106              xcoef = x(acol)
107              tmpsum = tmpsum + acoef*xcoef
108          enddo
109          y(i) = tmpsum
110      enddo
```



- ▶ Performs Mat/Vec product of sparse matrix
- ▶ Matrices are stored in a row-compressed format
- ▶ Parallelism per-row will vary, but is generally not very large

PARALLELIZED SPMV

```
106 !$acc parallel loop present(arrow_offsets,acols,acoefs) &
107 !$acc& private(row_start,row_end,acol,acoef,xcoef) &
108 !$acc& reduction(+:tmpsum)
109 do i=1,a%num_rows
110     tmpsum = 0.0d0
111     row_start = arrow_offsets(i)
112     row_end   = arrow_offsets(i+1)-1
113     do j=row_start,row_end
114         acol = acols(j)
115         acoef = acoefs(j)
116         xcoef = x(acol)
117         tmpsum = tmpsum + acoef*xcoef
118     enddo
119     y(i) = tmpsum
120 enddo
```

- ▶ Data already on device
- ▶ Compiler has vectorized the loop at 113 and selected a vector length of 256
- ▶ Total application speed-up (including other accelerated routines):
1.08X

OPENACC: 3 LEVELS OF PARALLELISM



- **Vector** threads work in lockstep (SIMD/SIMT parallelism)
- **Workers** compute a vector
- **Gangs** have 1 or more workers and share resources (such as cache, the streaming multiprocessor, etc.)
- Multiple gangs work independently of each other

OPENACC GANG, WORKER, VECTOR CLAUSES


- ▶ **gang**, **worker**, and **vector** can be added to a loop clause
- ▶ A parallel region can only specify one of each gang, worker, vector
- ▶ Control the size using the following clauses on the parallel region
 - ▶ **num_gangs(n)**, **num_workers(n)**, **vector_length(n)**

```
#pragma acc kernels loop gang
for (int i = 0; i < n; ++i)
    #pragma acc loop vector(128)
    for (int j = 0; j < n; ++j)
        ...
```

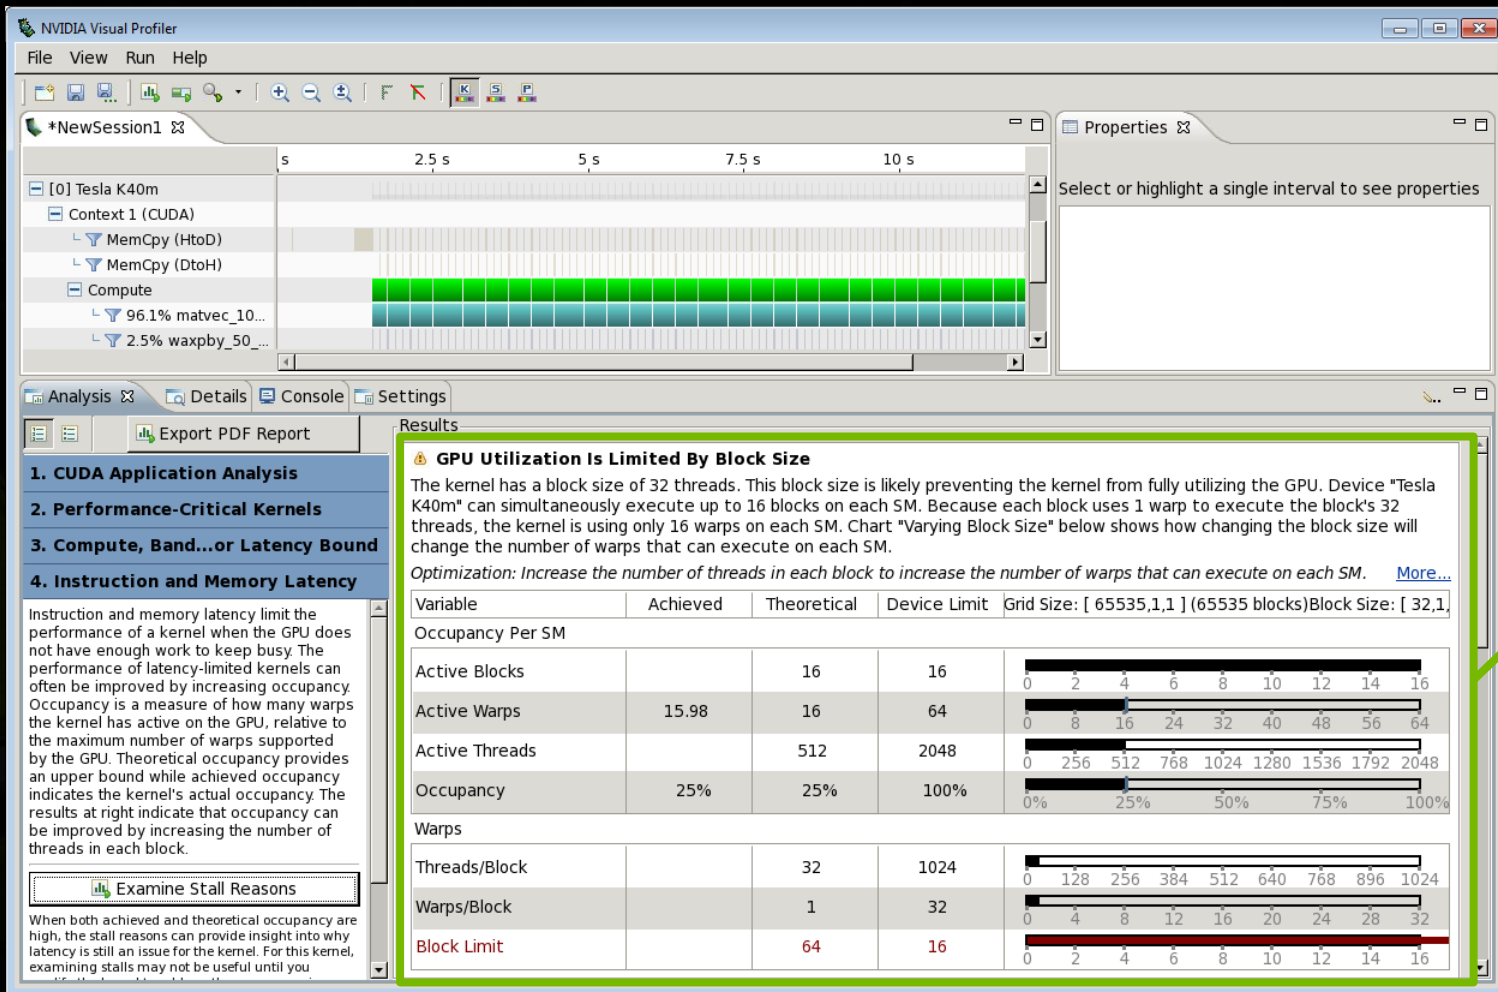
```
#pragma acc parallel vector_length(128)
#pragma acc loop gang
for (int i = 0; i < n; ++i)
    #pragma acc loop vector
    for (int j = 0; j < n; ++j)
        ...
```

OPTIMIZED SPMV VECTOR LENGTH

```
106 !$acc parallel loop present(arrow_offsets,acols,acoefs) &
107 !$acc& private(row_start,row_end,acol,acoef,xcoef) &
108 !$acc& vector_length(32)
109 do i=1,a%num_rows
110     tmpsum = 0.0d0
111     row_start = arrow_offsets(i)
112     row_end   = arrow_offsets(i+1)-1
113     !$acc loop vector reduction(+:tmpsum)
114     do j=row_start,row_end
115         acol = acols(j)
116         acoef = acoefs(j)
117         xcoef = x(acol)
118         tmpsum = tmpsum + acoef*xcoef
119     enddo
120     y(i) = tmpsum
121 enddo
```



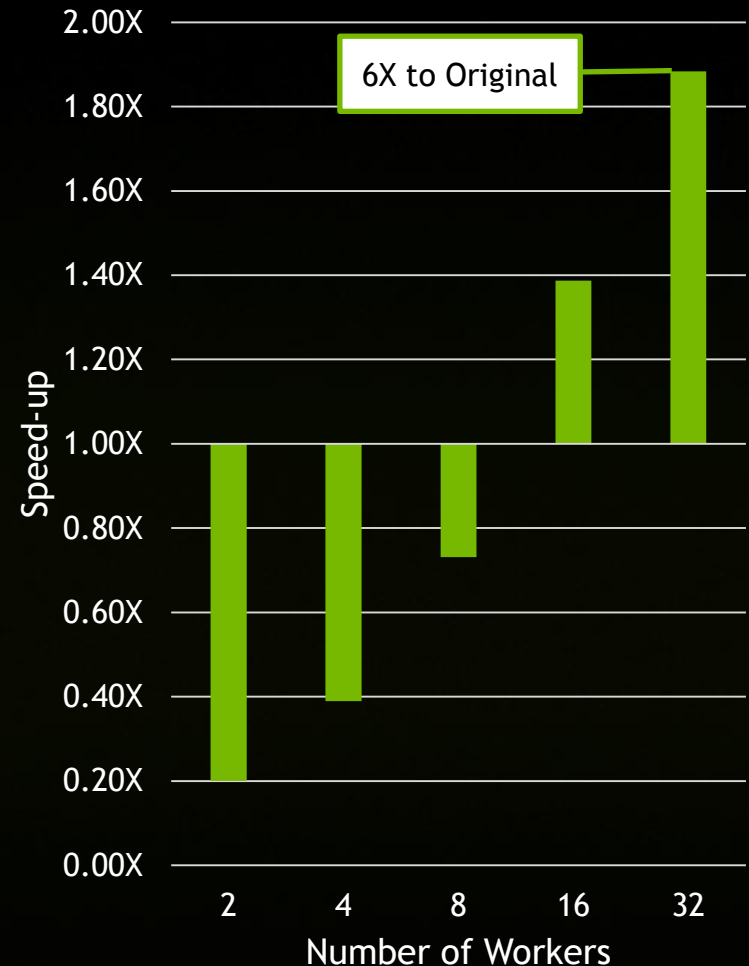
PERFORMANCE LIMITER: OCCUPANCY



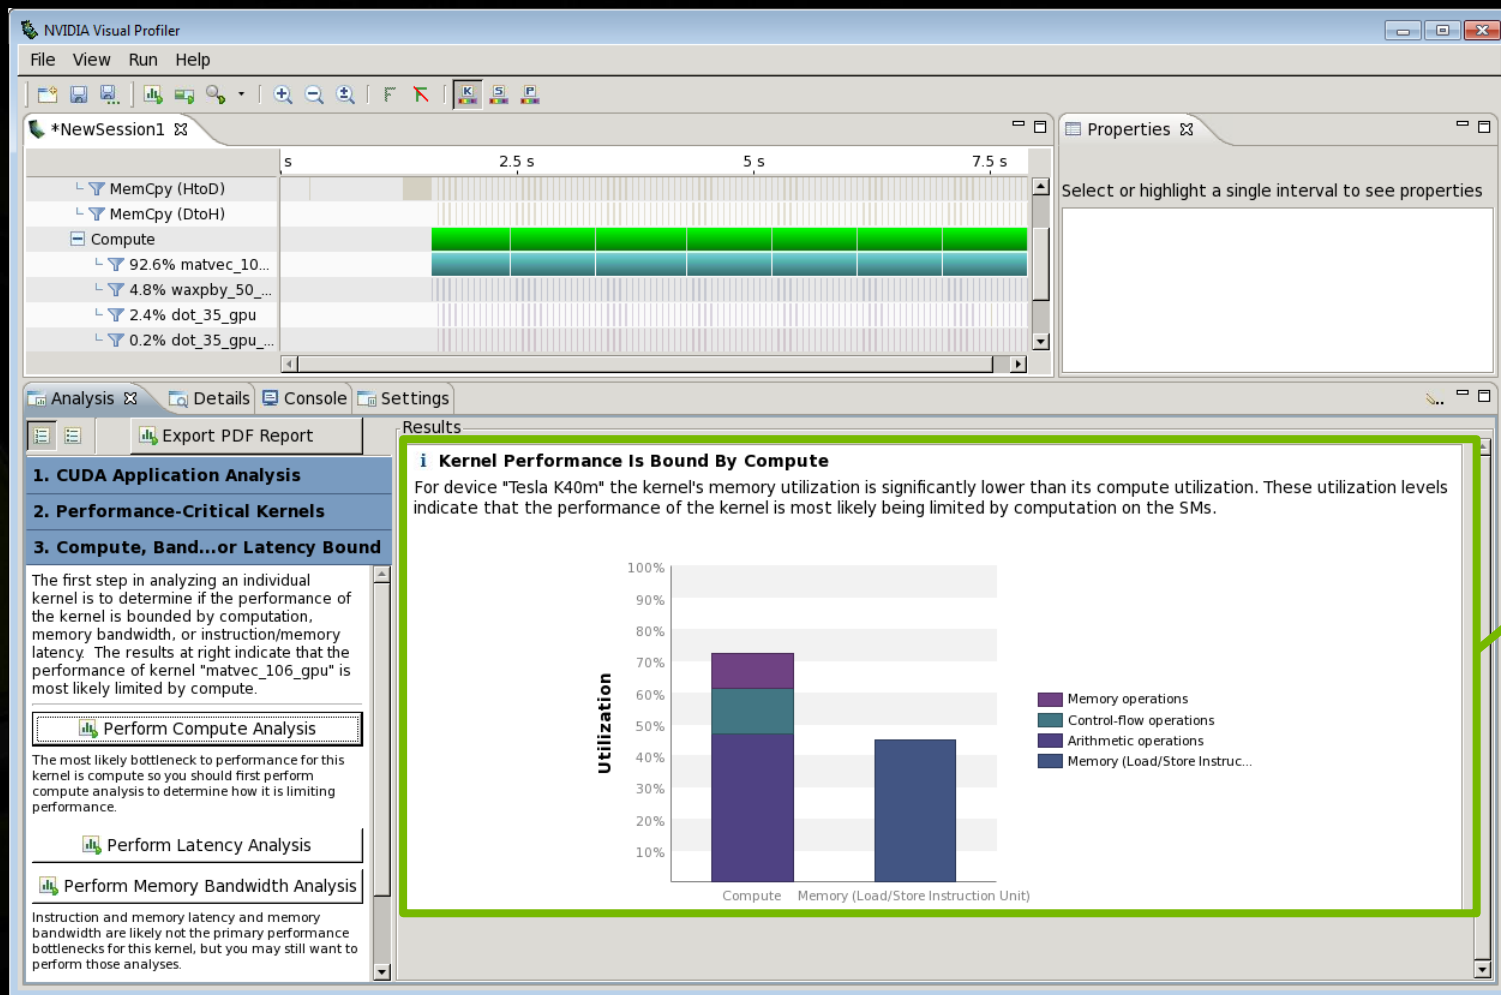
We need more threads!

INCREASED PARALLELISM WITH WORKERS

```
106 !$acc parallel loop present(arrow_offsets,acols,acoefs) &
107 !$acc& private(row_start,row_end,acol,acoef,xcoef) &
108 !$acc& gang worker vector_length(32) num_workers(32)
109 do i=1,a%num_rows
110     tmpsum = 0.0d0
111     row_start = arrow_offsets(i)
112     row_end   = arrow_offsets(i+1)-1
113     !$acc loop vector reduction(+:tmpsum)
114     do j=row_start,row_end
115         acol = acols(j)
116         acoef = acoefs(j)
117         xcoef = x(acol)
118         tmpsum = tmpsum + acoef*xcoef
119     enddo
120     y(i) = tmpsum
121 enddo
```

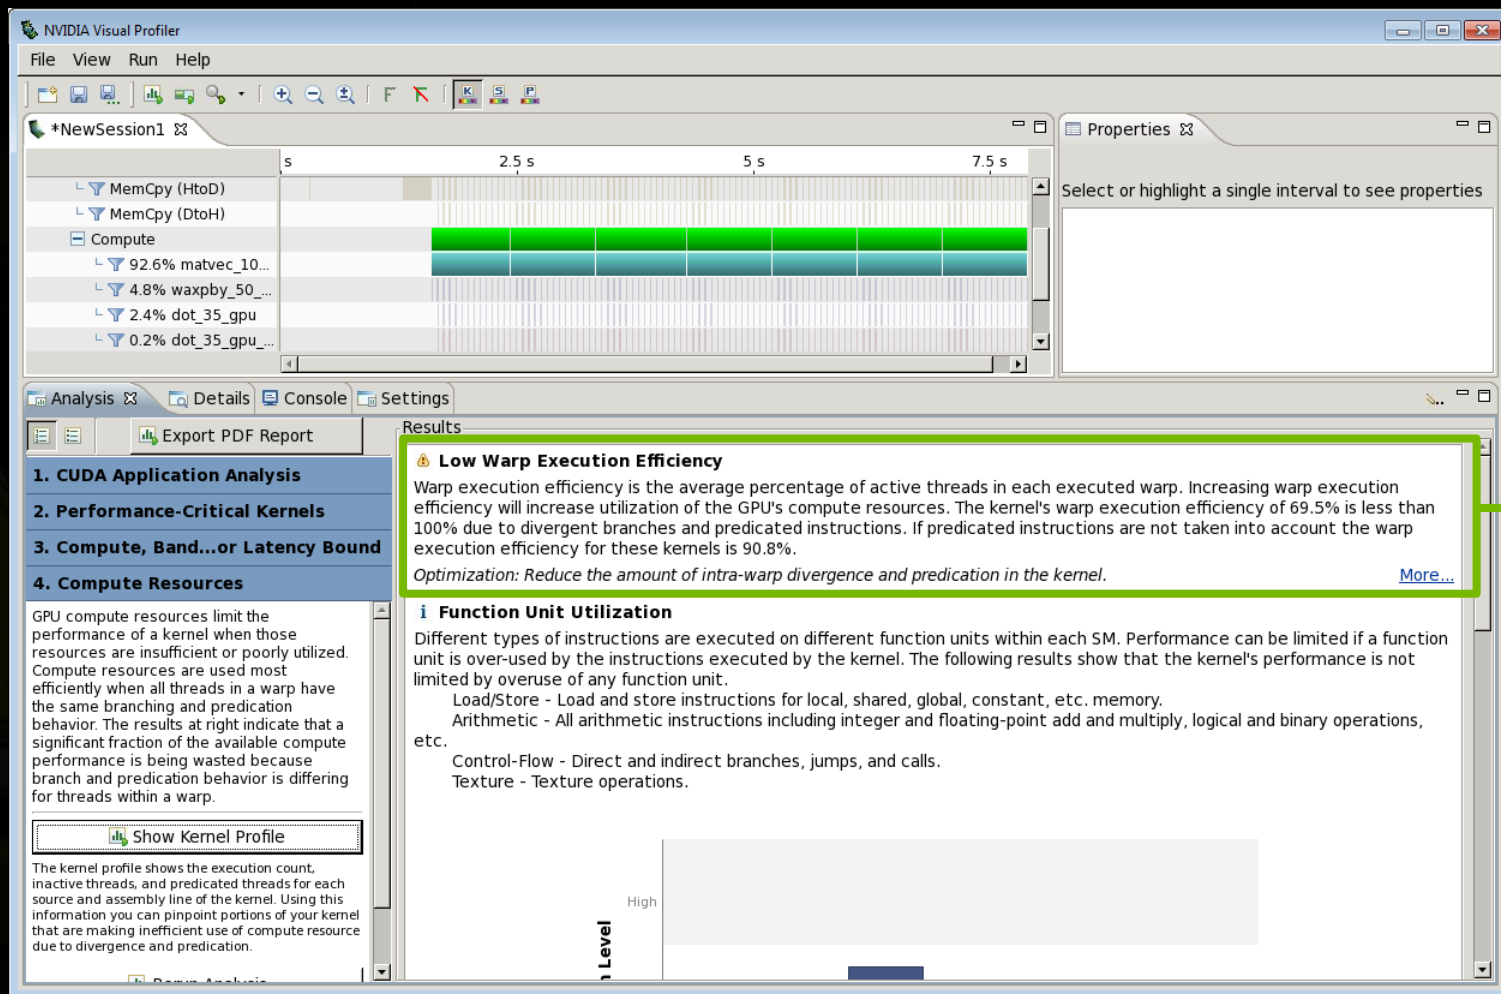


PERFORMANCE LIMITER: COMPUTE



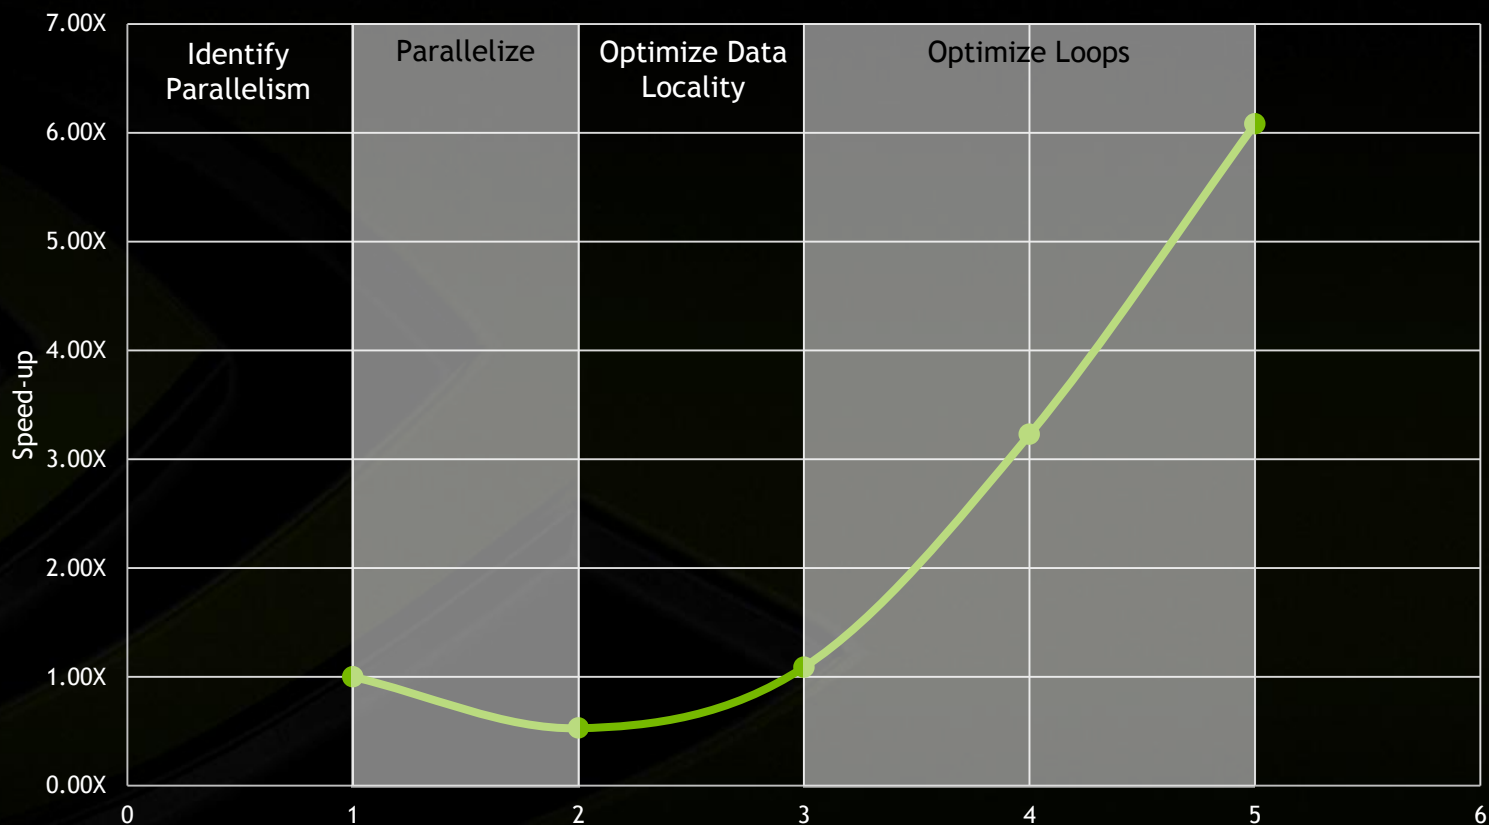
Now we're
compute bound

PERFORMANCE LIMITER: PARALLELISM



Really, we're limited by parallelism per-row.

SPEED-UP STEP BY STEP



OPENACC COLLAPSE CLAUSE

collapse(n): Transform the following n tightly nested loops into one, flattened loop.

- Useful when individual loops lack sufficient parallelism or more than 3 loops are nested (gang/worker/vector)

```
#pragma acc parallel
#pragma acc loop collapse(2)
for(int i=0; i<N; i++)
    for(int j=0; j<N; j++)
        ...
```




```
#pragma acc parallel
#pragma acc loop
for(int ij=0; ij<N*N; ij++)
    ...
```



Loops must be tightly nested

NEW CASE STUDY: MANDELBROT SET

- ▶ Application generates the image to the right.
- ▶ Each pixel in the image can be independently calculated.
- ▶ Skills Used:
 - ▶ Parallel Loop
 - ▶ Data Region
 - ▶ Update Directive
 - ▶ Asynchronous Pipelining



MANDELBROT CODE

```
// Calculate value for a pixel
unsigned char mandelbrot(int Px, int Py) {
    double x0=xmin+Px*dx;   double y0=ymin+Py*dy;
    double x=0.0;   double y=0.0;
    for(int i=0;x*x+y*y<4.0 && i<MAX_ITERS;i++) {
        double xtemp=x*x-y*y+x0;
        y=2*x*y+y0;
        x=xtemp;
    }
    return (double)MAX_COLOR*i/MAX_ITERS;
}

// Used in main()
for(int y=0;y<HEIGHT;y++) {
    for(int x=0;x<WIDTH;x++) {
        image[y*WIDTH+x]=mandelbrot(x,y);
    }
}
```

The mandelbrot() function calculates the color for each pixel.

Within main() there is a doubly-nested loop that calculates each pixel independently.



ROUTINES

OPENACC ROUTINE DIRECTIVE

Specifies that the compiler should generate a device copy of the function/subroutine and what type of parallelism the routine contains.

Clauses:

- ▶ **gang/worker/vector/seq**
 - ▶ Specifies the level of parallelism contained in the routine.
- ▶ **bind**
 - ▶ Specifies an optional name for the routine, also supplied at call-site
- ▶ **no_host**
 - ▶ The routine will only be used on the device
- ▶ **device_type**
 - ▶ Specialize this routine for a particular device type.

MANDELBROT: ROUTINE DIRECTIVE

```
// mandelbrot.h  
  
#pragma acc routine seq  
unsigned char mandelbrot(int Px, int Py);  
  
// Used in main()  
#pragma acc parallel loop  
for(int y=0;y<HEIGHT;y++) {  
    for(int x=0;x<WIDTH;x++) {  
        image[y*WIDTH+x]=mandelbrot(x,y);  
    }  
}
```

- ▶ At function source:
 - ▶ Function needs to be built for the GPU.
 - ▶ It will be called by each thread (sequentially)
- ▶ At call the compiler needs to know:
 - ▶ Function will be available on the GPU
 - ▶ It is a sequential routine

OPENACC ROUTINE: FORTRAN

```
module mandelbrot_mod
  implicit none
  integer, parameter :: HEIGHT=16384
  integer, parameter :: WIDTH=16384
  integer, parameter :: MAXCOLORS = 255

contains
  real(8) function mandlebrot(px,py)
    implicit none
    !$acc routine(mandlebrot) seq
    ...
  end function mandlebrot

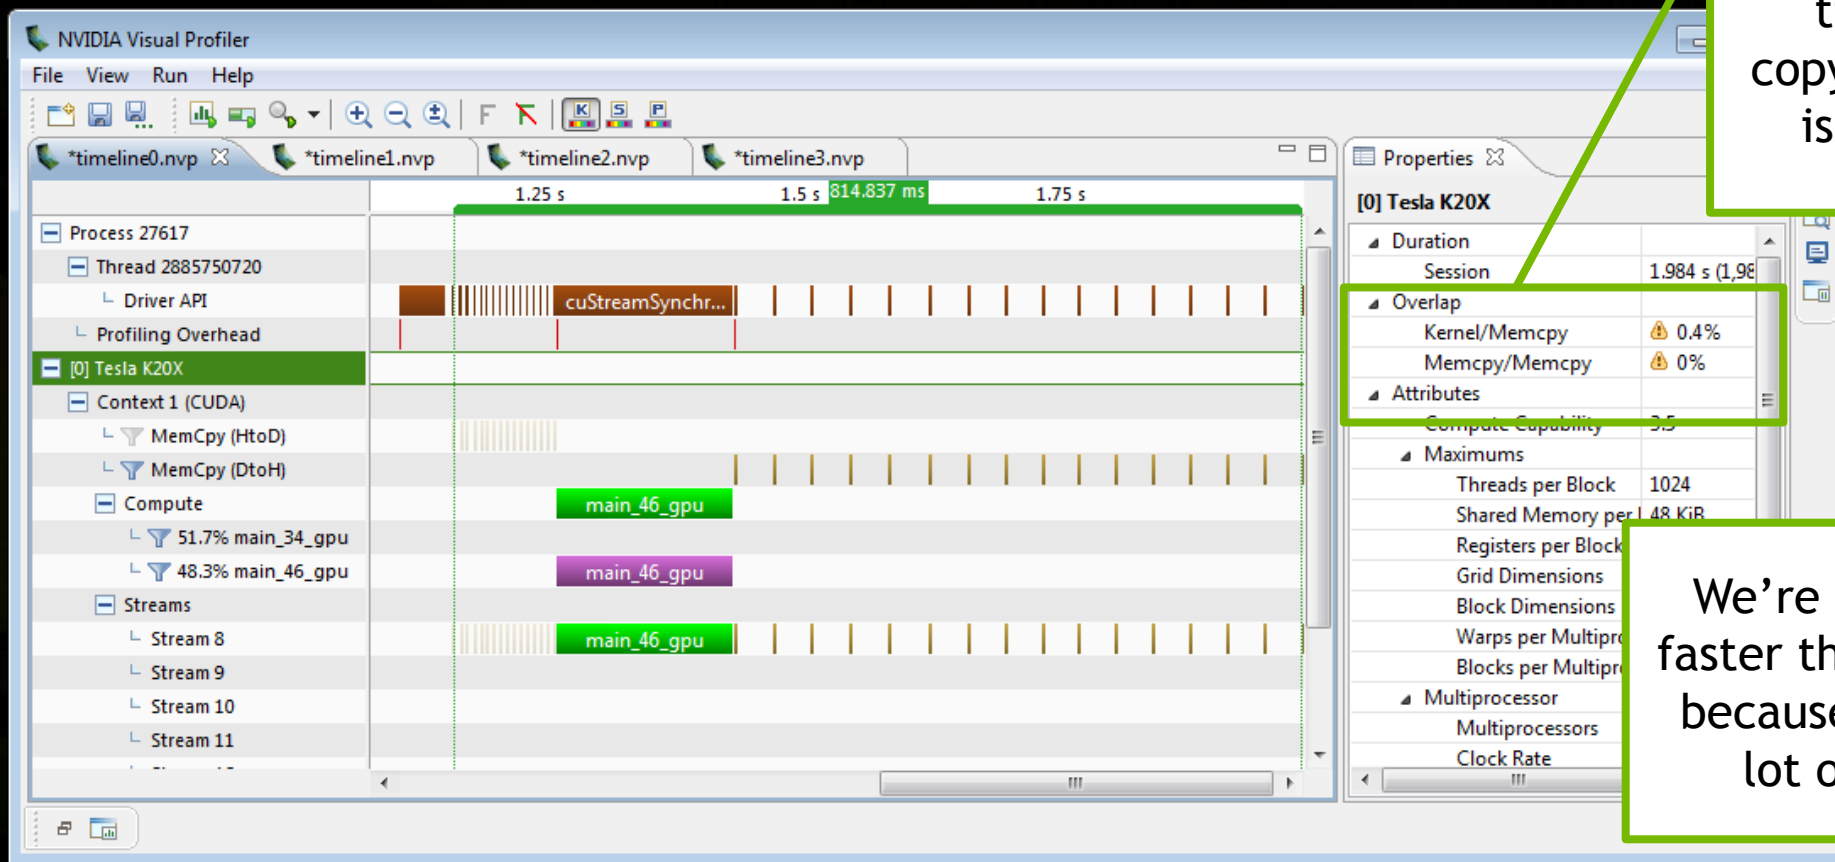
end module mandelbrot_mod
```

The **routine** directive may appear in a Fortran function or subroutine definition, or in an interface block.

The save attribute is not supported.

Nested acc routines require the routine directive within each nested routine.

BASLINE PROFILE



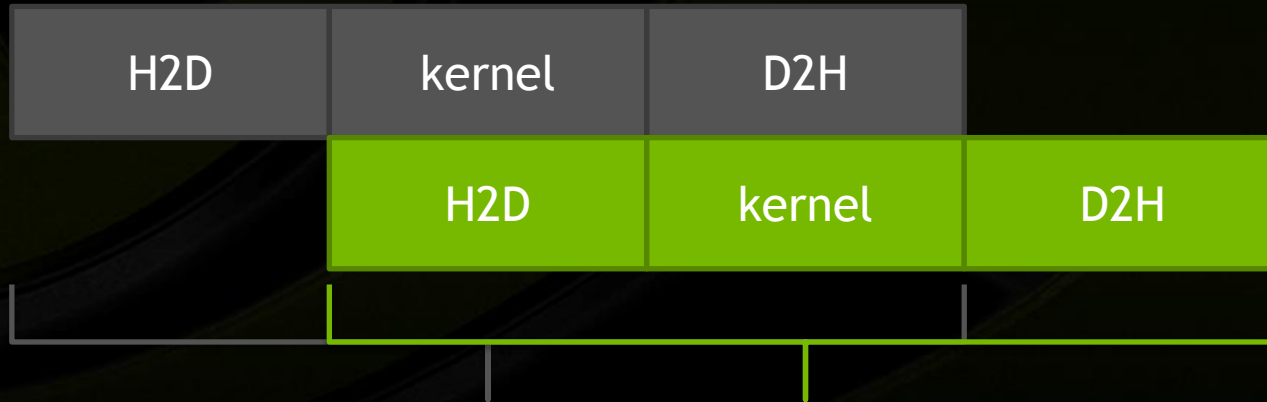
Roughly 25% of our time is spent copying, none of it is overlapped.

We're still much faster than the CPU because there's a lot of work.

PIPELINING DATA TRANSFERS



Two Independent Operations Serialized




Overlapping Copying and Computation

NOTE: In real applications, your boxes will not be so evenly sized.

PIPELINING MANDELBROT SET

- ▶ We only have 1 kernel, so there's nothing to overlap.
- ▶ Since each pixel is independent, computation can be broken up
- ▶ Steps
 1. Break up computation into blocks along rows.
 2. Break up copies according to blocks
 3. Make both computation and copies asynchronous

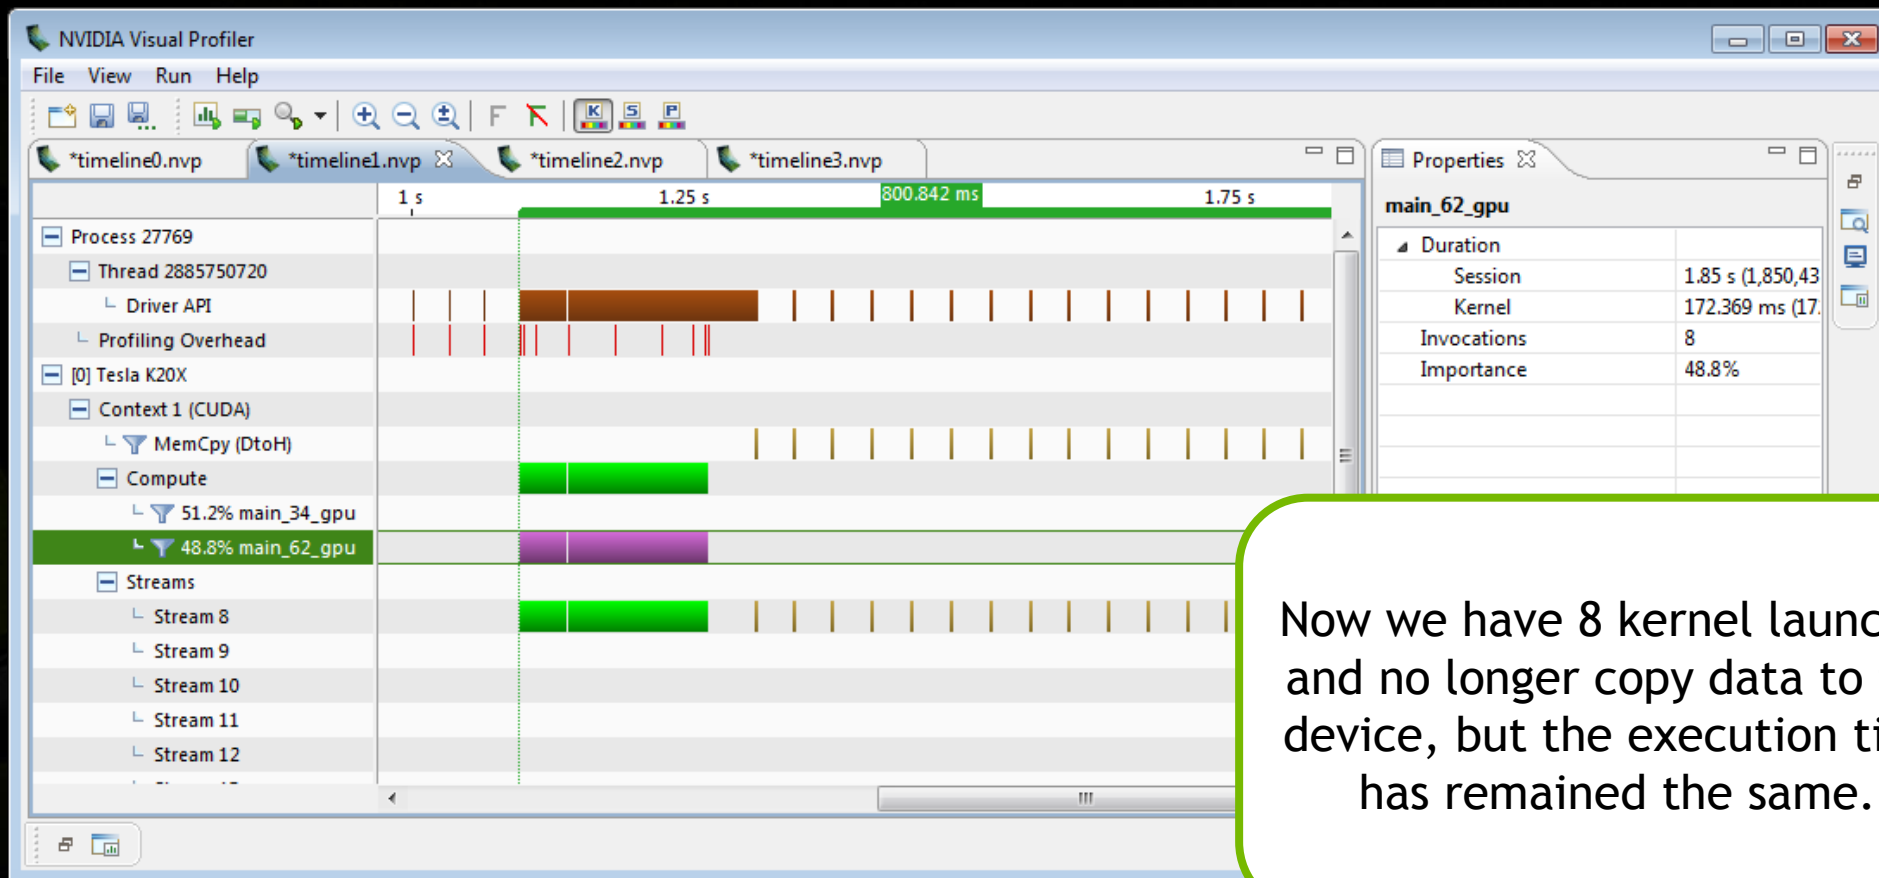


STEP 1: BLOCKING COMPUTATION

```
24  numblocks = ( argc > 1 ) ? atoi(argv[1]) : 8;
25  blocksize = HEIGHT / numblocks;
26  printf("numblocks: %d, blocksize: %d\n",
numblocks, blocksize);
27
28  #pragma acc data copyout(image[:bytes])
29  for(int block=0; block < numblocks; block++)
30  {
31      int ystart = block * blocksize;
32      int yend   = ystart + blocksize;
33  #pragma acc parallel loop
34      for(int y=ystart;y<yend;y++) {
35          for(int x=0;x<WIDTH;x++) {
36              image[y*WIDTH+x]=mandelbrot(x,y);
37          }
38      }
39  }
```

- ▶ Add a loop over blocks
- ▶ Modify the existing row loop to only work within blocks
- ▶ Add data region around blocking loop to leave data local to the device.
- ▶ Check for correct results.
- ▶ NOTE: We don't need to copy in the array, so make it an explicit copyout.

BLOCKING TIMELINE



Now we have 8 kernel launches and no longer copy data to the device, but the execution time has remained the same.




UPDATE DIRECTIVE

OPENACC DATA REGIONS REVIEW



```
28 #pragma acc data copyout(image[:bytes])
29   for(int block=0; block < numblocks; block++)
30   {
31       int ystart = block * blocksize;
32       int yend   = ystart + blocksize;
33   #pragma acc parallel loop
34       for(int y=ystart;y<yend;y++) {
35           for(int x=0;x<WIDTH;x++) {
36               image[y*WIDTH+x]=mandelbrot(x,y) ;
37           }
38       }
39   }
```

A large green curly bracket on the right side of the slide, spanning from line 28 to line 39 of the code block.

Data is shared
within this
region.

OPENACC UPDATE DIRECTIVE

Programmer specifies an array (or part of an array) that should be refreshed within a data region.

```
do_something_on_device()
```

```
!$acc update self(a)
```

```
do_something_on_host()
```

```
!$acc update device(a)
```



Copy “a” from GPU to
CPU



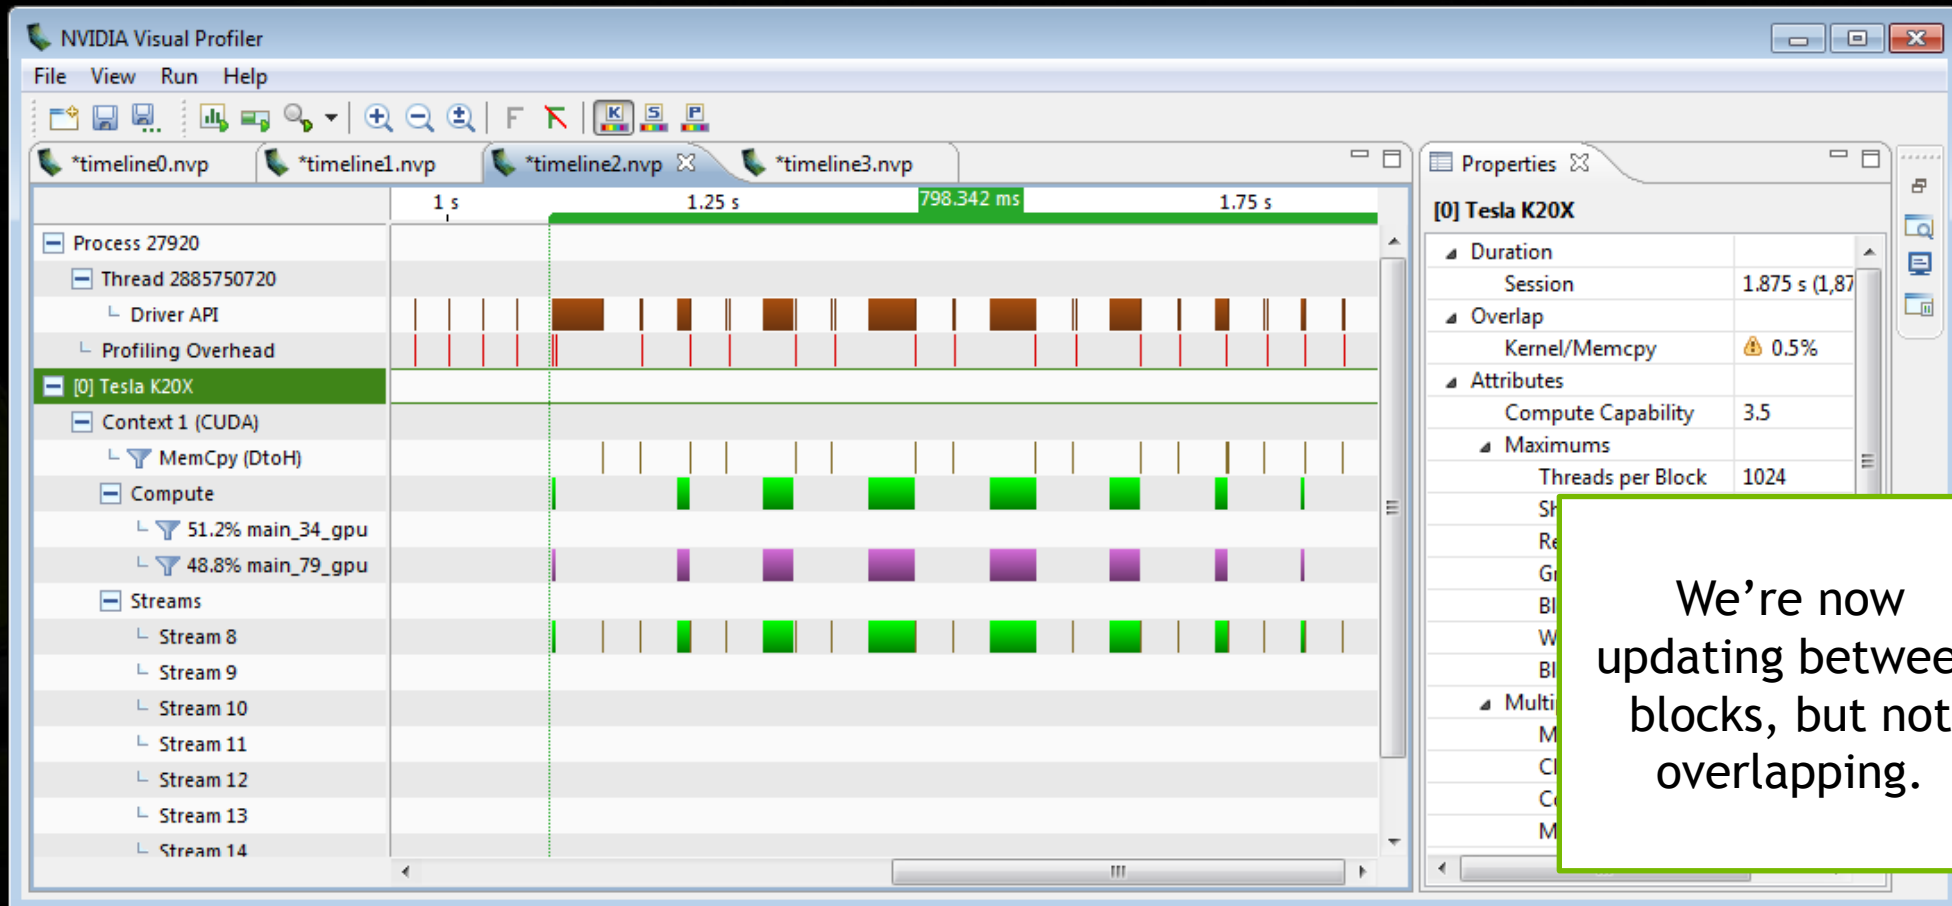
Copy “a” from CPU to
GPU

STEP 2: COPY BY BLOCK

```
28 #pragma acc data create(image[:bytes])
29   for(int block=0; block < numblocks; block++)
30   {
31       int ystart = block * blocksize;
32       int yend   = ystart + blocksize;
33   #pragma acc parallel loop
34       for(int y=ystart;y<yend;y++) {
35           for(int x=0;x<WIDTH;x++) {
36               image[y*WIDTH+x]=mandelbrot(x,y) ;
37           }
38       }
39   #pragma acc update
self(image[ystart*WIDTH:WIDTH*blocksize])
40   }
```

- ▶ Change the data region to only create the array on the GPU
- ▶ Use an update directive to copy individual blocks back to the host when complete
- ▶ Check for correct results.

TIMELINE: UPDATING BY BLOCKS





ASYNCHRONOUS PROGRAMMING

OPENACC ASYNC AND WAIT

async(n): launches work asynchronously in queue *n*

wait(n): blocks host until all operations in queue *n* have completed

Can significantly reduce launch latency, enables pipelining and concurrent operations

```
#pragma acc parallel loop async(1)
...
#pragma acc parallel loop async(1)
for(int i=0; i<N; i++)
    ...
#pragma acc wait(1)
for(int i=0; i<N; i++)
```

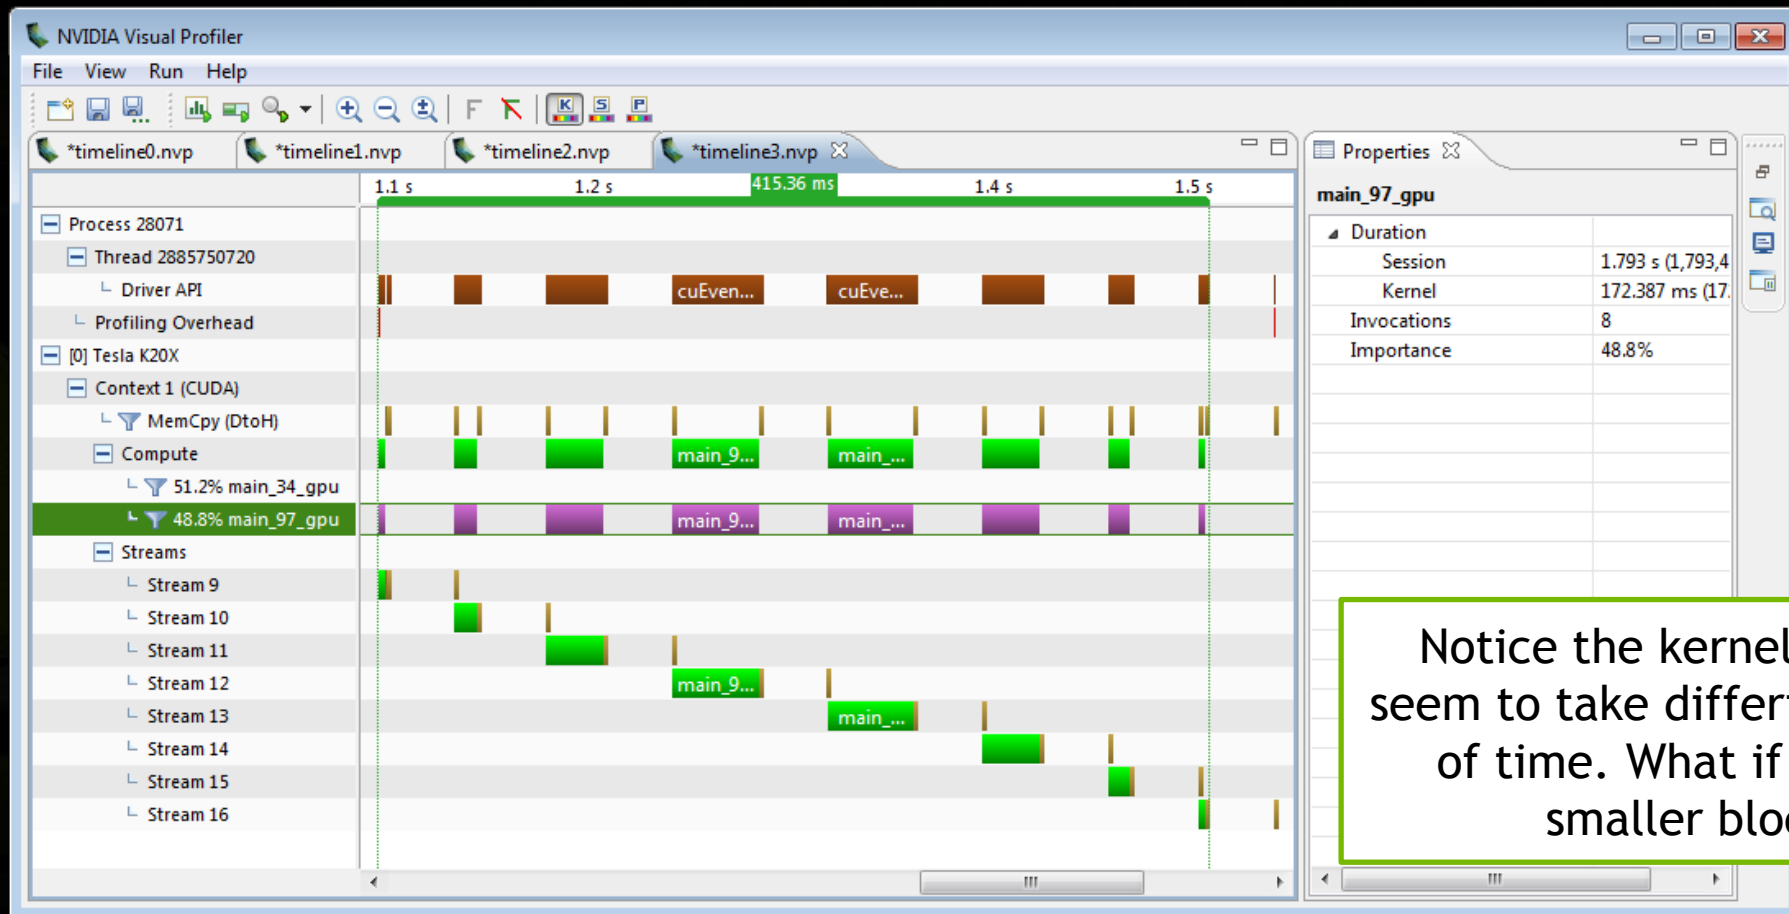
If *n* is not specified, *async* will go into a default queue and *wait* will wait all previously queued work.

STEP 3: GO ASYNCHRONOUS

```
31 #pragma acc data create(image[:bytes])
32   for(int block=0; block < numblocks; block++)
33   {
34       int ystart = block * blocksize;
35       int yend   = ystart + blocksize;
36 #pragma acc parallel loop async(block)
37       for(int y=ystart;y<yend;y++) {
38           for(int x=0;x<WIDTH;x++) {
39               image[y*WIDTH+x]=mandelbrot(x,y) ;
40           }
41       }
42 #pragma acc update
self(image[ystart*WIDTH:WIDTH*blocksize])
async(block)
43   }
44 #pragma acc wait
```


- ▶ Make each parallel region asynchronous by placing in different queues.
- ▶ Make each update asynchronous by placing in same stream as the parallel region on which it depends
- ▶ Synchronize for all to complete.
- ▶ Check for correct results.

TIMELINE: PIPELINING




Notice the kernel launches seem to take differing amounts of time. What if we tried smaller blocks?

VARYING THE NUMBER OF BLOCKS



SPEED-UP STEP BY STEP



ASYNCHRONOUS TIPS

- ▶ Reuse streams, they're expensive to create
 - ▶ Pre-create them
 - ▶ Consider `async(block%2)` to re-use just 2 streams
- ▶ Don't forget to `wait`
- ▶ Test with 1 stream first



MULTI-GPU PROGRAMMING

MULTI-GPU OPENACC

`acc_set_device_num(number, device_type)`

- ▶ Selects the device to use for all regions that follow

`acc_get_num_devices(device_type)`

- ▶ Queries how many devices are available of a given type
- ▶ Most often, one will set a device number once per CPU thread

MULTI-GPU MANDELBROT



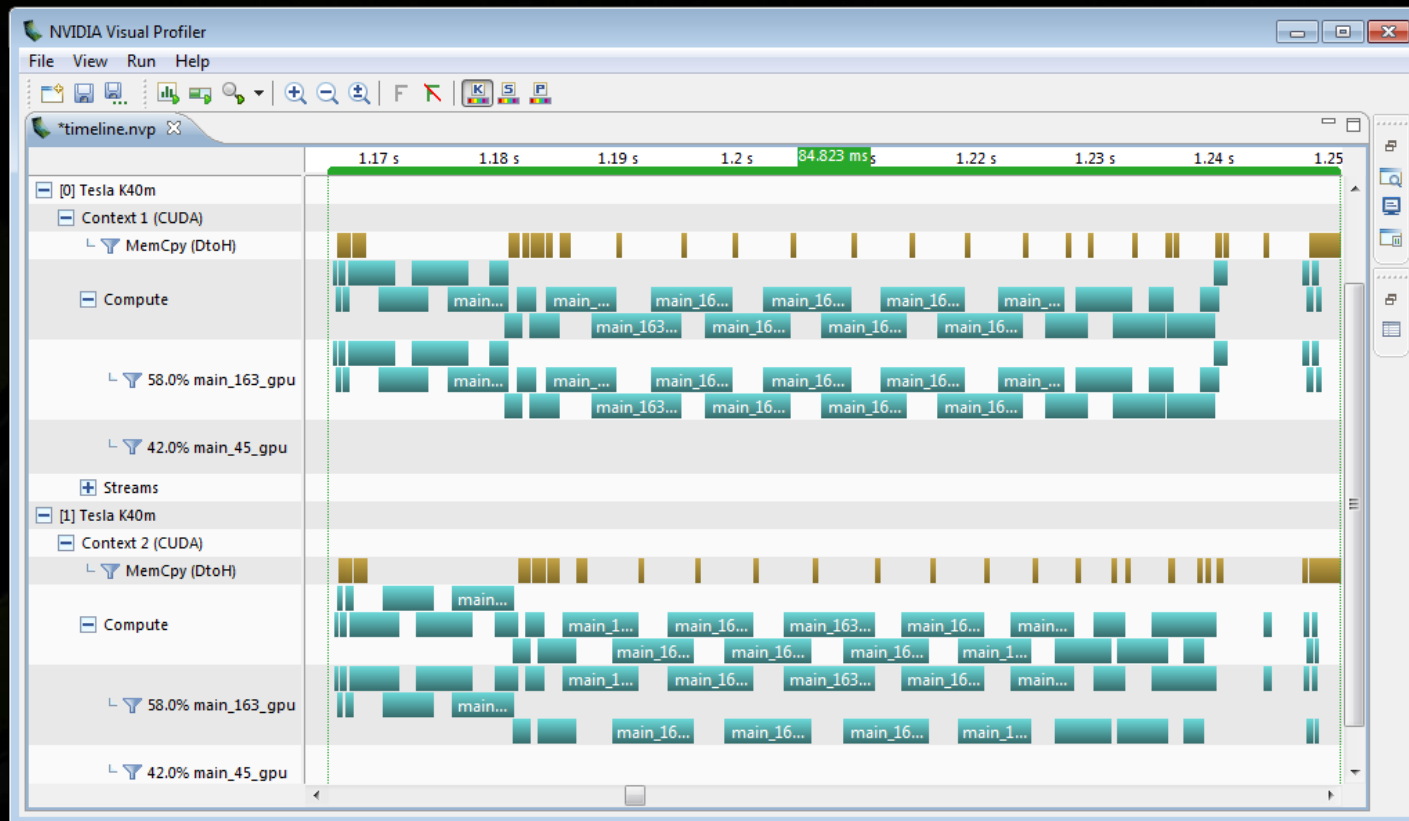
```
for (int gpu=0; gpu < 2 ; gpu ++)  
{  
    acc_set_device_num(gpu,acc_device_nvidia);  
#pragma acc enter data create(image[:bytes])  
}  
  
for(int block=0; block < numblocks; block++)  
{  
    int ystart = block * blocksize;  
    int yend   = ystart + blocksize;  
    acc_set_device_num(block%2,acc_device_nvidia);  
#pragma acc parallel loop async(block)  
    for(int y=ystart;y<yend;y++) {  
        for(int x=0;x<WIDTH;x++) {  
            image[y*WIDTH+x]=mandelbrot(x,y);  
        }  
    }  
#pragma acc update self(image[ystart*WIDTH:WIDTH*blocksize]) async(block)  
}  
for (int gpu=0; gpu < 2 ; gpu ++)  
{  
    acc_set_device_num(gpu,acc_device_nvidia);  
#pragma acc wait  
#pragma acc exit data delete(image)  
}
```

Allocate space on each device

Alternate devices per block

Clean up the devices

MULTI-GPU MANDELBROT PROFILE



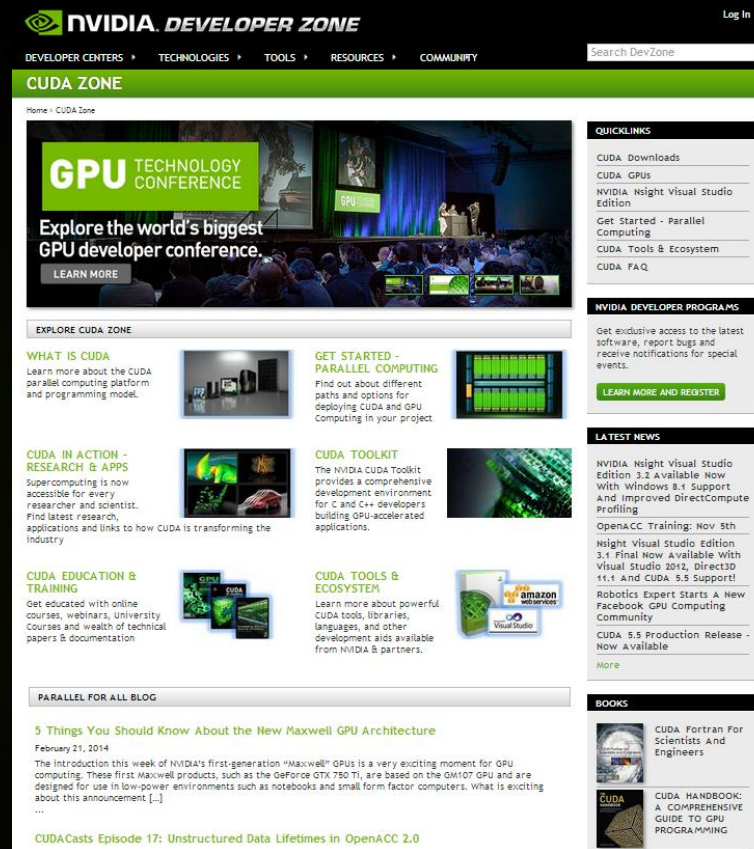


OPENACC INTEROPERABILITY

OPENACC INTEROPERABILITY

OpenACC plays well with others.

- ▶ Add CUDA or accelerated libraries to an OpenACC application
- ▶ Add OpenACC to an existing accelerated application
- ▶ Share data between OpenACC and CUDA



The screenshot shows the NVIDIA Developer Zone website. The header includes the NVIDIA logo, the text "NVIDIA DEVELOPER ZONE", a "Log In" link, and a search bar labeled "Search DevZone". Below the header is a green navigation bar with links for "DEVELOPER CENTERS", "TECHNOLOGIES", "TOOLS", "RESOURCES", and "COMMUNITY". The main content area is titled "CUDA ZONE" and features a large banner for the "GPU TECHNOLOGY CONFERENCE" with the text "Explore the world's biggest GPU developer conference." and a "LEARN MORE" button. Below the banner are several sections: "EXPLORE CUDA ZONE" with sub-sections like "WHAT IS CUDA", "GET STARTED - PARALLEL COMPUTING", "CUDA IN ACTION - RESEARCH & APPS", "CUDA TOOLKIT", "CUDA EDUCATION & TRAINING", and "CUDA TOOLS & ECOSYSTEM"; "QUICKLINKS" with links to "CUDA Downloads", "CUDA GPU", "NVIDIA Night Visual Studio Edition", "Get Started - Parallel Computing", "CUDA Tools & Ecosystem", and "CUDA FAQ"; "NVIDIA DEVELOPER PROGRAMS" with a "LEARN MORE AND REGISTER" button; "LATEST NEWS" with articles about "NVIDIA Night Visual Studio Edition 3.2 Available Now", "OpenACC Training: Nov 5th", "NVIDIA Night Visual Studio Edition 3.1 Final Now Available", and "CUDA 5.5 Production Release - Now Available"; and "BOOKS" with links to "CUDA Fortran For Scientists and Engineers" and "CUDA HANDBOOK: A COMPREHENSIVE GUIDE TO GPU PROGRAMMING". At the bottom, there is a "PARALLEL FOR ALL BLOG" section with an article titled "5 Things You Should Know About the New Maxwell GPU Architecture" and a "CUDAcasts Episode 17: Unstructured Data Lifetimes in OpenACC 2.0" link.

OPENACC & CUDA STREAMS

OpenACC *suggests* two functions for interoperating with CUDA streams:

- ▶ `void* acc_get_cuda_stream(int async);`
- ▶ `int acc_set_cuda_stream(int async, void* stream);`

OPENACC HOST_DATA DIRECTIVE

Exposes the *device* address of particular objects to the *host* code.

```
#pragma acc data copy(x,y)
{
    // x and y are host pointers
    #pragma acc host_data use_device(x,y)
    {
        // x and y are device pointers
    }
    // x and y are host pointers
}
```

} X and Y are device
pointers here

HOST_DATA EXAMPLE



OpenACC Main

```
program main
  integer, parameter :: N = 2**20
  real, dimension(N) :: X, Y
  real                :: A = 2.0

  !$acc data
  ! Initialize X and Y
  ...

  !$acc host_data use_device(x,y)
  call saxpy(n, a, x, y)
  !$acc end host_data
  !$acc end data

end program
```

- It's possible to interoperate from C/C++ or Fortran.
- OpenACC manages the data and passes device pointers to CUDA.

CUDA C Kernel & Wrapper

```
__global__
void saxpy_kernel(int n, float a,
                  float *x, float *y)
{
  int i = blockIdx.x*blockDim.x + threadIdx.x;
  if (i < n) y[i] = a*x[i] + y[i];
}

void saxpy(int n, float a, float *dx, float *dy)
{
  // Launch CUDA Kernel
  saxpy_kernel<<<4096,256>>>(N, 2.0, dx, dy);
}
```

- CUDA kernel launch wrapped in function expecting device arrays.
- Kernel is launch with arrays passed from OpenACC in main.

CUBLAS LIBRARY & OPENACC



OpenACC can interface with existing GPU-optimized libraries (from C/C++ or Fortran).

This includes...

- CUBLAS
- Libsci_acc
- CUFFT
- MAGMA
- CULA
- Thrust
- ...

OpenACC Main Calling CUBLAS

```
int N = 1<<20;
float *x, *y
// Allocate & Initialize x & y
...

cublasInit();

#pragma acc data copyin(x[0:N]) copy(y[0:N])
{
    #pragma acc host_data use_device(x,y)
    {
        // Perform SAXPY on 1M elements
        cublasSaxpy(N, 2.0, x, 1, y, 1);
    }
}

cublasShutdown();
```

OPENACC DEVICEPTR

The **deviceptr** clause informs the compiler that an object is already on the device, so no translation is necessary.

- ▶ Valid for **parallel**, **kernels**, and **data**

```
cudaMallocManaged((void*)&x,(size_t)n*sizeof(float));  
cudaMallocManaged((void*)&y,(size_t)n*sizeof(float));
```

```
#pragma acc parallel loop deviceptr(x,y)  
for(int i=0; i<n ; i++)  
{  
    y(i) = a*x(i)+y(i)  
}
```

Do not translate x
and y, they are
already on the
device.

DEVICEPTR EXAMPLE

OpenACC Kernels

```
void saxpy(int n, float a, float * restrict
x, float * restrict y)
{
    #pragma acc kernels deviceptr(x[0:n],y[0:n])
    {
        for(int i=0; i<n; i++)
        {
            y[i] += 2.0*x[i];
        }
    }
}
```

By passing a device pointer to an OpenACC region, it's possible to add OpenACC to an existing CUDA code.

CUDA C Main

```
int main(int argc, char **argv)
{
    float *x, *y, tmp;
    int n = 1<<20, i;

    cudaMalloc((void*)&x,(size_t)n*sizeof(float));
    cudaMalloc((void*)&y,(size_t)n*sizeof(float));

    ...

    saxpy(n, 2.0, x, y);
    cudaMemcpy(&tmp,y,(size_t)sizeof(float),
               cudaMemcpyDeviceToHost);

    return 0;
}
```

Memory is managed via standard CUDA calls.

OPENACC & THRUST

Thrust (thrust.github.io) is a STL-like library for C++ on accelerators.

- High-level interface
- Host/Device container classes
- Common parallel algorithms

It's possible to cast Thrust vectors to device pointers for use with OpenACC

```
void saxpy(int n, float a, float * restrict
x, float * restrict y)
{
#pragma acc kernels deviceptr(x[0:n],y[0:n])
{
    for(int i=0; i<n; i++)
    {
        y[i] += 2.0*x[i];
    }
}
}
```

```
int N = 1<<20;
thrust::host_vector<float> x(N), y(N);
for(int i=0; i<N; i++)
{
    x[i] = 1.0f;
    y[i] = 0.0f;
}

// Copy to Device
thrust::device_vector<float> d_x = x;
thrust::device_vector<float> d_y = y;

saxpy(N, 2.0, d_x.data().get(),
      d_y.data().get());

// Copy back to host
y = d_y;
```

CUDA DEVICE ROUTINES AND OPENACC

```
extern "C" __device__ void  
f1dev( float* a, float* b, int i ){  
    a[i] = .... b[i] .... ;  
}
```

Even CUDA `__device__` functions
can be called from OpenACC if
declared with `acc routine`.

```
#pragma acc routine seq  
extern "C" void f1dev( float*,  
float* int );  
...  
#pragma acc parallel loop \  
    present( a[0:n], b[0:n] )  
for( int i = 0; i < n; ++i )  
{  
    f1dev( a, b, i );  
}
```

OPENACC ACC_MAP_DATA FUNCTION

The `acc_map_data` (`acc_unmap_data`) maps (unmaps) an existing device allocation to an OpenACC variable.

```
cudaMalloc((void*)&x_d,(size_t)n*sizeof(float));  
acc_map_data(x, x_d, n*sizeof(float));  
cudaMalloc((void*)&y_d,(size_t)n*sizeof(float));  
acc_map_data(y, y_d, n*sizeof(float));
```

```
#pragma acc parallel loop  
for(int i=0; i<n ; i++)  
{  
    y(i) = a*x(i)+y(i)  
}
```

Allocate device
arrays with CUDA
and *map* to
OpenACC

Here x and y will
reuse the memory
of x_d and y_d



ATOMIC DIRECTIVE

OPENACC ATOMIC DIRECTIVE


atomic: subsequent block of code is performed atomically with respect to other threads on the accelerator

Clauses: **read, write, update, capture**

```
#pragma acc parallel loop
for(int i=0; i<N; i++) {
    #pragma acc atomic update
    a[i%100]++;
}
```

OPENACC ATOMIC: HISTOGRAM

```
19  #pragma acc data copyin(a[0:N]) copyout(h[0:HN])
20  for(int it=0;it<ITERS;it++)
21  {
22      #pragma acc parallel loop
23      for(int i=0;i<HN;i++)
24          h[i]=0;
25
26      #pragma acc parallel loop
27      for(int i=0;i<N;i++) {
28          #pragma acc atomic
29          h[a[i]]+=1;
30      }
31  }
```





MISC. ADVICE AND TECHNIQUES

WRITE PARALLELIZABLE LOOPS

Use countable loops
C99: while->for
Fortran: while->do

Avoid pointer
arithmetic (use
array syntax)

Write rectangular
loops (compiler
cannot parallelize
triangular loops)

```
bool found=false;
while(!found && i<N) {
    if(a[i]==val) {
        found=true
        loc=i;
    }
    i++;
}
```

```
bool found=false;
for(int i=0;i<N;i++) {
    if(a[i]==val) {
        found=true
        loc=i;
    }
}
```

```
for(int i=0;i<N;i++) {
    for(int j=i;j<N;j++) {
        sum+=A[i][j];
    }
}
```

```
for(int i=0;i<N;i++) {
    for(int j=0;j<N;j++) {
        if(j>=i)
            sum+=A[i][j];
    }
}
```

ALIASING CAN PREVENT PARALLELIZATION

23, Loop is parallelizable

Accelerator kernel generated

```
23, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
```

25, Complex loop carried dependence of 'b->' prevents parallelization

Loop carried dependence of 'a->' prevents parallelization

Loop carried backward dependence of 'a->' prevents vectorization

Accelerator scalar kernel generated

27, Complex loop carried dependence of 'a->' prevents parallelization

Loop carried dependence of 'b->' prevents parallelization

Loop carried backward dependence of 'b->' prevents vectorization

Accelerator scalar kernel generated

C99: RESTRICT KEYWORD

- ▶ Declaration of intent given by the programmer to the compiler

Applied to a pointer, e.g.

```
float *restrict ptr
```

Meaning: “for the lifetime of `ptr`, only it or a value directly derived from it (such as `ptr + 1`) will be used to access the object to which it points”*

- ▶ Parallelizing compilers often require `restrict` to determine independence
 - ▶ Otherwise the compiler can’t parallelize loops that access `ptr`
 - ▶ Note: if programmer violates the declaration, behavior is undefined



```
float restrict *ptr  
float *restrict ptr
```

<http://en.wikipedia.org/wiki/Restrict>

INLINING

- ▶ When possible aggressively inline functions/routines
 - ▶ This is especially important for inner loop calculations
 - ▶ Inlined routines frequently perform better than acc routines because the compiler has more information.

```
#pragma acc routine seq
inline
int IDX(int row, int col, int LDA) {
    return row*LDA+col;
}
```

KERNEL FUSION

- ▶ Kernel calls are expensive
 - ▶ Each call can take over 10us in order to launch
 - ▶ It is often a good idea to combine loops of same trip counts containing very few lines of code
- ▶ Kernel Fusion (i.e. Loop fusion)
 - ▶ Join nearby kernels into a single kernel

```
#pragma acc parallel loop
  for (int i = 0; i < n; ++i) {
    a[i]=0;
  }
#pragma acc parallel loop
  for (int i = 0; i < n; ++i) {
    b[i]=0;
  }
```



```
#pragma acc parallel loop
  for (int i = 0; i < n; ++i) {
    a[i]=0;
    b[i]=0;
  }
```

LOOP FISSION

- ▶ Loops that are exceptionally long may result in kernels that are resource-bound, resulting in low GPU occupancy.
- ▶ This is particularly true for outer parallel loops containing nested loops
- ▶ Caution: This may introduce temporaries.


```
#pragma acc parallel loop
for (int j = 0; j < m; ++j ) {
    for (int i = 0; i < n; ++i) {
        a[i]=0;
    }
    for (int i = 0; i < n; ++i) {
        b[i]=0;
    }
}
```



```
#pragma acc parallel loop
for (int j = 0; j < m; ++j )
    for (int i = 0; i < n; ++i) {
        a[i]=0;
    }
#pragma acc parallel loop
for (int j = 0; j < m; ++j )
    for (int i = 0; i < n; ++i) {
        b[i]=0;
    }
```

MEMORY COALESCING

- ▶ *Coalesced* access:
 - ▶ A group of 32 contiguous threads (“warp”) accessing adjacent words
 - ▶ Few transactions and high utilization
- ▶ *Uncoalesced* access:
 - ▶ A warp of 32 threads accessing scattered words
 - ▶ Many transactions and low utilization
- ▶ For best performance the **vector** loop should access memory **contiguously (stride-1)**



COMPLEX DATA LAYOUTS

- ▶ OpenACC works best with flat arrays
- ▶ Some compilers handle complex types (structs, classes, derived types) better than others
 - ▶ Doesn't always work, particularly if members are dynamically allocated
 - ▶ Work around: Use local pointers to struct members (C99 & Fortran)

May work

```
#pragma acc parallel loop \  
        copy(a, a.data[0:a.N])  
for (i=0;i<a.N;i++)  
    a.data[i]=0;
```

Generally Works

```
int N=a.N;  
float *data=a.data;  
#pragma acc parallel loop \  
        copy(data[0:N])  
for (i=0;i<N;i++)  
    data[i]=0;
```


OPENACC INDEPENDENT CLAUSE

Specifies that loop iterations are data independent. This overrides any compiler dependency analysis. This is implied for *parallel loop*.

```
#pragma acc kernels
```

```
{  
#pragma acc loop independent  
for(int i=0; i<N; i++)  
{  
    a[i] = 0.0;  
    b[i] = 1.0;  
    c[i] = 2.0;  
}  
#pragma acc loop independent  
for(int i=0; i<N; i++)  
{  
    a(i) = b(i) + c(i)  
}  
}
```

kernel 1

kernel 2

Informs the compiler that both loops are safe to parallelize so it will generate both kernels.

OPENACC DEBUGGING

- ▶ Most OpenACC directives accept an `if(condition)` clause

```
#pragma acc update self(A) if(debug)
```

```
#pragma acc parallel loop if(!debug)
```

```
[...]
```

```
#pragma acc update device(A) if(debug)
```

- ▶ Use `default(none)` to force explicit data directives

```
#pragma acc data copy(...) create(...) default(none)
```



QUESTIONS & DISCUSSION